

```
from google.colab import files
files.upload()
```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

```
'kaggle.json': {'username': 'adityamunaga', 'key': '0c876f34h72h7247d4h7fh76a97ha9h'}}
```

```
!mkdir ~/.kaggle/
```

```
!cp kaggle.json ~/.kaggle/
```

```
!chmod 600 ~/.kaggle/kaggle.json
```

```
!kaggle competitions download -c dogs-vs-cats
```

```
Downloading dogs-vs-cats.zip to /content
```

```
100% 810M/812M [00:28<00:00, 27.5MB/s]
```

```
100% 812M/812M [00:28<00:00, 29.5MB/s]
```

```
!unzip -qq dogs-vs-cats.zip
```

```
!unzip -qq train.zip
```

```
import os, shutil, pathlib
```

```
original_dir = pathlib.Path("train")
```

```
new_base_dir = pathlib.Path("cats_vs_dogs_small_1")
```

```
def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)
```

```
make_subset("train", start_index=0, end_index=1000)
```

```
make_subset("validation", start_index=1000, end_index=1500)
```

```
make_subset("test", start_index=1500, end_index=2000)
```

```
from tensorflow import keras
```

```
from tensorflow.keras import layers
```

```
inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

```
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 180, 180, 3)]	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896

max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
conv2d_2 (Conv2D)	(None, 41, 41, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	0
conv2d_3 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 256)	0
conv2d_4 (Conv2D)	(None, 7, 7, 256)	590080
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 1)	12545

```

=====
Total params: 991041 (3.78 MB)
Trainable params: 991041 (3.78 MB)
Non-trainable params: 0 (0.00 Byte)

```

---

```

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

```

from tensorflow.keras.utils import image_dataset_from_directory

```

```

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)

```

```

Found 2000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.

```

```

from tensorflow import keras

```

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch1.x",
        save_best_only=True,
        monitor="val_loss"
    )
]

```

```

history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

```

Epoch 1/10
63/63 [=====] - 168s 3s/step - loss: 0.6985 - accuracy: 0.5030 - val_loss: 0.6924 - val_accuracy: 0.5160
Epoch 2/10
63/63 [=====] - 191s 3s/step - loss: 0.6961 - accuracy: 0.5185 - val_loss: 0.6856 - val_accuracy: 0.6300
Epoch 3/10
63/63 [=====] - 161s 3s/step - loss: 0.6913 - accuracy: 0.5580 - val_loss: 0.6768 - val_accuracy: 0.5560
Epoch 4/10
63/63 [=====] - 181s 3s/step - loss: 0.6617 - accuracy: 0.6060 - val_loss: 0.6251 - val_accuracy: 0.6440
Epoch 5/10
63/63 [=====] - 181s 3s/step - loss: 0.6182 - accuracy: 0.6515 - val_loss: 0.7889 - val_accuracy: 0.5620
Epoch 6/10
63/63 [=====] - 161s 3s/step - loss: 0.5904 - accuracy: 0.6915 - val_loss: 0.6178 - val_accuracy: 0.6480

```

```

Epoch 7/10
63/63 [=====] - 156s 2s/step - loss: 0.5599 - accuracy: 0.7050 - val_loss: 0.6965 - val_accuracy: 0.6460
Epoch 8/10
63/63 [=====] - 179s 3s/step - loss: 0.5259 - accuracy: 0.7390 - val_loss: 0.6757 - val_accuracy: 0.6710
Epoch 9/10
63/63 [=====] - 165s 3s/step - loss: 0.4827 - accuracy: 0.7695 - val_loss: 0.6239 - val_accuracy: 0.7130
Epoch 10/10
63/63 [=====] - 182s 3s/step - loss: 0.4600 - accuracy: 0.7800 - val_loss: 0.5699 - val_accuracy: 0.7350

```



It can be shown that as the number of epochs increases, accuracy increases.

Accuracy=77.2% Val\_acc=67.7% test accuracy=63.6%

```

import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()

```



```
test_model = keras.models.load_model("convnet_from_scratch1.x")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
```

```
32/32 [=====] - 20s 627ms/step - loss: 0.5584 - accuracy: 0.7330
Test accuracy: 0.733
```

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)
```

```
inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

```
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

```
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation1.x",
        save_best_only=True,
        monitor="val_loss")
]
```

```
history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset,
    callbacks=callbacks)
```

```
Epoch 1/10
63/63 [=====] - 185s 3s/step - loss: 0.7005 - accuracy: 0.5190 - val_loss: 0.6927 - val_accuracy: 0.5000
Epoch 2/10
63/63 [=====] - 195s 3s/step - loss: 0.6929 - accuracy: 0.5265 - val_loss: 0.7032 - val_accuracy: 0.5000
Epoch 3/10
63/63 [=====] - 197s 3s/step - loss: 0.6936 - accuracy: 0.5450 - val_loss: 0.6878 - val_accuracy: 0.5050
Epoch 4/10
63/63 [=====] - 178s 3s/step - loss: 0.6956 - accuracy: 0.5430 - val_loss: 0.6946 - val_accuracy: 0.5090
Epoch 5/10
63/63 [=====] - 194s 3s/step - loss: 0.6737 - accuracy: 0.5985 - val_loss: 0.6526 - val_accuracy: 0.6460
Epoch 6/10
63/63 [=====] - 171s 3s/step - loss: 0.6689 - accuracy: 0.6015 - val_loss: 0.6510 - val_accuracy: 0.6020
Epoch 7/10
63/63 [=====] - 175s 3s/step - loss: 0.6366 - accuracy: 0.6470 - val_loss: 1.1360 - val_accuracy: 0.5010
Epoch 8/10
63/63 [=====] - 171s 3s/step - loss: 0.6149 - accuracy: 0.6650 - val_loss: 1.1188 - val_accuracy: 0.5110
Epoch 9/10
63/63 [=====] - 172s 3s/step - loss: 0.6079 - accuracy: 0.6565 - val_loss: 0.5953 - val_accuracy: 0.6600
Epoch 10/10
63/63 [=====] - 169s 3s/step - loss: 0.5946 - accuracy: 0.6825 - val_loss: 0.6078 - val_accuracy: 0.6580
```

```
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation1.x")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
```

```
32/32 [=====] - 19s 579ms/step - loss: 0.6203 - accuracy: 0.6740
Test accuracy: 0.674
```

```

import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small_2")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir, exist_ok=True)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=1500)
make_subset("validation", start_index=1500, end_index=2000)
make_subset("test", start_index=2000, end_index=2500)

from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)

    Found 3000 files belonging to 2 classes.
    Found 1000 files belonging to 2 classes.
    Found 1000 files belonging to 2 classes.

inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch2.x",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset,
    callbacks=callbacks)

Epoch 1/10
94/94 [=====] - 229s 2s/step - loss: 0.6978 - accuracy: 0.5293 - val_loss: 0.6835 - val_accuracy: 0.5450
Epoch 2/10
94/94 [=====] - 225s 2s/step - loss: 0.6782 - accuracy: 0.5913 - val_loss: 0.6605 - val_accuracy: 0.6280
Epoch 3/10
94/94 [=====] - 225s 2s/step - loss: 0.6375 - accuracy: 0.6577 - val_loss: 0.6477 - val_accuracy: 0.6400

```

```

Epoch 4/10
94/94 [=====] - 249s 3s/step - loss: 0.6086 - accuracy: 0.6680 - val_loss: 0.6404 - val_accuracy: 0.6220
Epoch 5/10
94/94 [=====] - 224s 2s/step - loss: 0.5624 - accuracy: 0.7040 - val_loss: 0.7732 - val_accuracy: 0.5960
Epoch 6/10
94/94 [=====] - 225s 2s/step - loss: 0.5234 - accuracy: 0.7410 - val_loss: 0.5908 - val_accuracy: 0.6910
Epoch 7/10
94/94 [=====] - 223s 2s/step - loss: 0.4714 - accuracy: 0.7777 - val_loss: 0.5460 - val_accuracy: 0.7240
Epoch 8/10
94/94 [=====] - 226s 2s/step - loss: 0.4135 - accuracy: 0.8093 - val_loss: 0.6275 - val_accuracy: 0.7230
Epoch 9/10
94/94 [=====] - 245s 3s/step - loss: 0.3662 - accuracy: 0.8390 - val_loss: 0.5625 - val_accuracy: 0.7350
Epoch 10/10
94/94 [=====] - 224s 2s/step - loss: 0.3049 - accuracy: 0.8720 - val_loss: 0.6284 - val_accuracy: 0.7160

test_model = keras.models.load_model(
    "convnet_from_scratch2.x")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 19s 583ms/step - loss: 0.5138 - accuracy: 0.7490
Test accuracy: 0.749

data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="adam",
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation2.x",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset,
    callbacks=callbacks)

Epoch 1/10
94/94 [=====] - 249s 3s/step - loss: 0.6924 - accuracy: 0.5263 - val_loss: 0.7081 - val_accuracy: 0.5400
Epoch 2/10
94/94 [=====] - 262s 3s/step - loss: 0.6800 - accuracy: 0.5670 - val_loss: 0.6654 - val_accuracy: 0.6150
Epoch 3/10
94/94 [=====] - 239s 3s/step - loss: 0.6567 - accuracy: 0.6160 - val_loss: 0.6599 - val_accuracy: 0.6280
Epoch 4/10
94/94 [=====] - 241s 3s/step - loss: 0.6521 - accuracy: 0.6243 - val_loss: 0.6278 - val_accuracy: 0.6620
Epoch 5/10
94/94 [=====] - 238s 3s/step - loss: 0.6334 - accuracy: 0.6547 - val_loss: 0.6521 - val_accuracy: 0.6100
Epoch 6/10
94/94 [=====] - 240s 3s/step - loss: 0.6209 - accuracy: 0.6590 - val_loss: 0.6373 - val_accuracy: 0.6230

```

```

Epoch 7/10
94/94 [=====] - 246s 3s/step - loss: 0.5840 - accuracy: 0.6870 - val_loss: 0.5703 - val_accuracy: 0.6920
Epoch 8/10
94/94 [=====] - 242s 3s/step - loss: 0.5841 - accuracy: 0.6830 - val_loss: 0.6220 - val_accuracy: 0.6550
Epoch 9/10
94/94 [=====] - 244s 3s/step - loss: 0.5527 - accuracy: 0.7227 - val_loss: 0.5654 - val_accuracy: 0.7050
Epoch 10/10
94/94 [=====] - 242s 3s/step - loss: 0.5464 - accuracy: 0.7197 - val_loss: 0.5550 - val_accuracy: 0.7200

```

```

test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation2.x")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

```

32/32 [=====] - 20s 620ms/step - loss: 0.5373 - accuracy: 0.7340
Test accuracy: 0.734

```

```
import os, shutil, pathlib
```

```

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small_3")

```

```

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir, exist_ok=True)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

```

```

make_subset("train", start_index=0, end_index=1500)
make_subset("validation", start_index=1500, end_index=2500)
make_subset("test", start_index=2500, end_index=3000)

```

```
from tensorflow.keras.utils import image_dataset_from_directory
```

```

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)

```

```

Found 3000 files belonging to 2 classes.
Found 2000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.

```

```

inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

```

```

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch3.x",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset,
    callbacks=callbacks)

Epoch 1/10
94/94 [=====] - 249s 3s/step - loss: 0.6994 - accuracy: 0.5160 - val_loss: 0.6895 - val_accuracy: 0.5220
Epoch 2/10
94/94 [=====] - 244s 3s/step - loss: 0.6874 - accuracy: 0.5627 - val_loss: 0.7314 - val_accuracy: 0.5005
Epoch 3/10
94/94 [=====] - 246s 3s/step - loss: 0.6567 - accuracy: 0.6133 - val_loss: 0.6334 - val_accuracy: 0.6595
Epoch 4/10
94/94 [=====] - 246s 3s/step - loss: 0.6124 - accuracy: 0.6650 - val_loss: 0.6216 - val_accuracy: 0.6475
Epoch 5/10
94/94 [=====] - 245s 3s/step - loss: 0.5933 - accuracy: 0.6907 - val_loss: 0.5856 - val_accuracy: 0.7010
Epoch 6/10
94/94 [=====] - 244s 3s/step - loss: 0.5427 - accuracy: 0.7253 - val_loss: 0.6702 - val_accuracy: 0.6860
Epoch 7/10
94/94 [=====] - 244s 3s/step - loss: 0.5017 - accuracy: 0.7550 - val_loss: 0.6098 - val_accuracy: 0.7130
Epoch 8/10
94/94 [=====] - 242s 3s/step - loss: 0.4589 - accuracy: 0.7877 - val_loss: 0.6190 - val_accuracy: 0.7145
Epoch 9/10
94/94 [=====] - 246s 3s/step - loss: 0.4087 - accuracy: 0.8193 - val_loss: 0.6753 - val_accuracy: 0.6815
Epoch 10/10
94/94 [=====] - 245s 3s/step - loss: 0.3539 - accuracy: 0.8377 - val_loss: 0.6167 - val_accuracy: 0.7460

test_model = keras.models.load_model(
    "convnet_from_scratch3.x")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 20s 598ms/step - loss: 0.6104 - accuracy: 0.6740
Test accuracy: 0.674

data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="adam",
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation3.x",
        save_best_only=True,
        monitor="val_loss")
]

```



```

history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset,
    callbacks=callbacks)

Epoch 1/10
94/94 [=====] - 312s 3s/step - loss: 0.6944 - accuracy: 0.5037 - val_loss: 0.6940 - val_accuracy: 0.5000
Epoch 2/10
94/94 [=====] - 264s 3s/step - loss: 0.6936 - accuracy: 0.5080 - val_loss: 0.6939 - val_accuracy: 0.5005
Epoch 3/10
94/94 [=====] - 267s 3s/step - loss: 0.6888 - accuracy: 0.5160 - val_loss: 0.6921 - val_accuracy: 0.5960
Epoch 4/10
94/94 [=====] - 262s 3s/step - loss: 0.6806 - accuracy: 0.5470 - val_loss: 0.6942 - val_accuracy: 0.5465
Epoch 5/10
94/94 [=====] - 262s 3s/step - loss: 0.6643 - accuracy: 0.6027 - val_loss: 0.6617 - val_accuracy: 0.5900
Epoch 6/10
94/94 [=====] - 263s 3s/step - loss: 0.6530 - accuracy: 0.6147 - val_loss: 0.6355 - val_accuracy: 0.6610
Epoch 7/10
94/94 [=====] - 262s 3s/step - loss: 0.6250 - accuracy: 0.6690 - val_loss: 0.6479 - val_accuracy: 0.6190
Epoch 8/10
94/94 [=====] - 263s 3s/step - loss: 0.6274 - accuracy: 0.6567 - val_loss: 0.6278 - val_accuracy: 0.6550
Epoch 9/10
94/94 [=====] - 264s 3s/step - loss: 0.6185 - accuracy: 0.6600 - val_loss: 0.6249 - val_accuracy: 0.6575
Epoch 10/10
94/94 [=====] - 265s 3s/step - loss: 0.6062 - accuracy: 0.6713 - val_loss: 0.5829 - val_accuracy: 0.6925

```

```

test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation3.x")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 22s 643ms/step - loss: 0.6131 - accuracy: 0.6760
Test accuracy: 0.676

```

```

conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False,
    input_shape=(180, 180, 3))
conv_base.summary()

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop\\_58889256/58889256](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop_58889256/58889256) [=====] - 3s 0us/step  
Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_7 (InputLayer)	[None, 180, 180, 3]	0
block1_conv1 (Conv2D)	(None, 180, 180, 64)	1792
block1_conv2 (Conv2D)	(None, 180, 180, 64)	36928
block1_pool (MaxPooling2D)	(None, 90, 90, 64)	0
block2_conv1 (Conv2D)	(None, 90, 90, 128)	73856
block2_conv2 (Conv2D)	(None, 90, 90, 128)	147584
block2_pool (MaxPooling2D)	(None, 45, 45, 128)	0
block3_conv1 (Conv2D)	(None, 45, 45, 256)	295168
block3_conv2 (Conv2D)	(None, 45, 45, 256)	590080
block3_conv3 (Conv2D)	(None, 45, 45, 256)	590080
block3_pool (MaxPooling2D)	(None, 22, 22, 256)	0
block4_conv1 (Conv2D)	(None, 22, 22, 512)	1180160
block4_conv2 (Conv2D)	(None, 22, 22, 512)	2359808
block4_conv3 (Conv2D)	(None, 22, 22, 512)	2359808
block4_pool (MaxPooling2D)	(None, 11, 11, 512)	0
block5_conv1 (Conv2D)	(None, 11, 11, 512)	2359808
block5_conv2 (Conv2D)	(None, 11, 11, 512)	2359808

```
block5_conv3 (Conv2D)      (None, 11, 11, 512)      2359808
```

```
block5_pool (MaxPooling2D) (None, 5, 5, 512)      0
```

```
=====
Total params: 14714688 (56.13 MB)
Trainable params: 14714688 (56.13 MB)
Non-trainable params: 0 (0.00 Byte)
```

---

```
import numpy as np
```

```
def get_features_and_labels(dataset):
    all_features = []
    all_labels = []
    for images, labels in dataset:
        preprocessed_images = keras.applications.vgg16.preprocess_input(images)
        features = conv_base.predict(preprocessed_images)
        all_features.append(features)
        all_labels.append(labels)
    return np.concatenate(all_features), np.concatenate(all_labels)
```

```
train_features, train_labels = get_features_and_labels(train_dataset)
val_features, val_labels = get_features_and_labels(validation_dataset)
test_features, test_labels = get_features_and_labels(test_dataset)
```

```
train_features.shape
```

```
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 13s 13s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 12s 12s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 12s 12s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 12s 12s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 12s 12s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 13s 13s/step
```

```

1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 10s 10s/step
1/1 [=====] - 11s 11s/step
1/1 [=====] - 11s 11s/step



inputs = keras.Input(shape=(5, 5, 512))
x = layers.Flatten()(inputs)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extractionPT1.x",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_features, train_labels,
    epochs=15,
    validation_data=(val_features, val_labels),
    callbacks=callbacks)

Epoch 1/15
94/94 [=====] - 16s 149ms/step - loss: 13.3925 - accuracy: 0.9357 - val_loss: 5.2545 - val_accuracy: 0.9625
Epoch 2/15
94/94 [=====] - 6s 60ms/step - loss: 4.1733 - accuracy: 0.9747 - val_loss: 9.2914 - val_accuracy: 0.9585
Epoch 3/15
94/94 [=====] - 6s 67ms/step - loss: 1.4125 - accuracy: 0.9887 - val_loss: 4.9386 - val_accuracy: 0.9710
Epoch 4/15
94/94 [=====] - 5s 51ms/step - loss: 0.9839 - accuracy: 0.9893 - val_loss: 5.6567 - val_accuracy: 0.9755
Epoch 5/15
94/94 [=====] - 7s 76ms/step - loss: 1.1170 - accuracy: 0.9923 - val_loss: 4.9057 - val_accuracy: 0.9760
Epoch 6/15
94/94 [=====] - 4s 46ms/step - loss: 0.5845 - accuracy: 0.9957 - val_loss: 5.6758 - val_accuracy: 0.9765
Epoch 7/15
94/94 [=====] - 4s 46ms/step - loss: 0.4703 - accuracy: 0.9970 - val_loss: 7.3636 - val_accuracy: 0.9695
Epoch 8/15
94/94 [=====] - 5s 53ms/step - loss: 0.4880 - accuracy: 0.9953 - val_loss: 6.0083 - val_accuracy: 0.9745
Epoch 9/15
94/94 [=====] - 4s 47ms/step - loss: 0.3857 - accuracy: 0.9953 - val_loss: 6.1635 - val_accuracy: 0.9750
Epoch 10/15
94/94 [=====] - 7s 73ms/step - loss: 0.2643 - accuracy: 0.9987 - val_loss: 6.1887 - val_accuracy: 0.9755
Epoch 11/15
94/94 [=====] - 5s 49ms/step - loss: 0.0613 - accuracy: 0.9993 - val_loss: 6.1906 - val_accuracy: 0.9755
Epoch 12/15
94/94 [=====] - 6s 60ms/step - loss: 0.4566 - accuracy: 0.9967 - val_loss: 5.9209 - val_accuracy: 0.9750
Epoch 13/15
94/94 [=====] - 5s 54ms/step - loss: 0.1049 - accuracy: 0.9987 - val_loss: 5.8666 - val_accuracy: 0.9775
Epoch 14/15
94/94 [=====] - 5s 49ms/step - loss: 0.1882 - accuracy: 0.9980 - val_loss: 11.3728 - val_accuracy: 0.9660
Epoch 15/15
94/94 [=====] - 5s 58ms/step - loss: 0.1420 - accuracy: 0.9980 - val_loss: 7.2477 - val_accuracy: 0.9720

import matplotlib.pyplot as plt
acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, "bo", label="Training accuracy")
plt.plot(epochs, val_acc, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()

```



```
conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False)
conv_base.trainable = False
```

```
conv_base.trainable = True
print("This is the number of trainable weights "
      "before freezing the conv base:", len(conv_base.trainable_weights))
```

```
conv_base.trainable = False
print("This is the number of trainable weights "
      "after freezing the conv base:", len(conv_base.trainable_weights))
```

```
This is the number of trainable weights before freezing the conv base: 26
This is the number of trainable weights after freezing the conv base: 0
```

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
```

```

x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extraction_with_data_augmentationPT2.x",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=5,
    validation_data=validation_dataset,
    callbacks=callbacks)

Epoch 1/5
94/94 [=====] - 1768s 19s/step - loss: 19.0464 - accuracy: 0.9070 - val_loss: 5.1251 - val_accuracy: 0.9680
Epoch 2/5
94/94 [=====] - 1686s 18s/step - loss: 6.5667 - accuracy: 0.9530 - val_loss: 4.1836 - val_accuracy: 0.9765
Epoch 3/5
94/94 [=====] - 1693s 18s/step - loss: 5.9281 - accuracy: 0.9577 - val_loss: 4.6219 - val_accuracy: 0.9740
Epoch 4/5
61/94 [=====>.....] - ETA: 5:53 - loss: 4.2165 - accuracy: 0.9652

```